AI / Data / ML / Cloud / Platform - Engineer | Developer | Data Scientist

Stijn M. W. Hering

LinkedIn
Amsterdam
www.stijnhering.nl
https://github.com/melchiorhering

Profile

Full-Stack Engineer specializing in Data, AI/ML, Cloud/DevOps, and Web Development. Outside of work-related engineering projects, I actively homelab, experiment with emerging tools, and tinker with new tech—continually expanding my skill set and putting practical innovations into practice

How would colleagues describe you?

Open, technically strong team player (DISC: Red/Yellow) who takes initiative and thrives in dynamic environments. I enjoy collaborating, solving complex problems, and driving innovative ideas into practical solutions.

What should colleagues know about you?

As a middle child in a family of 7, I grew up with strong negotiation skills, thick skin for constructive feedback.. and, of course, an open mind. Don't let my height (nearly 2m) fool you—I'm a friendly, reliable team player who loves collaborating on innovative solutions.

Achievements

Agentic Platform | Rabobank

- **Engineered a Cloud-Agnostic Agentic Platform:** Architected and deployed an open-source, self-hostable platform using Temporal and Langfuse, leveraging IaC, CI/CD, and robust engineering best practices.
- **Accelerated AI/Agentic Delivery:** Designed standardized agentic workflows and reusable project templates, applying software engineering standards to dramatically speed up development cycles.
- **Guardrails & Observability:** Implemented end-to-end evaluation, real-time monitoring, and security guardrails to ensure reliability and maintainability in production agent systems.

Used: AWS, Azure, Pulumi, Databricks, LangChain, Pydantic-AI, Temporal, Langfuse, Python, Typescript, React, AG-UI, Docker, GIT, FOSS.

Data/ML, DevOps & Cloud | TenneT

- **Engineered Legacy-to-Cloud Migration:** Contributed to the end-to-end migration of legacy on-premises Data and ML systems to a modern Azure & Databricks environment.
- **Streamlined Data & ML Deployments:** Established DevOps and CI/CD best practices, build the standardized project template to accelerate model and pipeline deployments on Databricks.
- **Enhanced Cloud Data Infrastructure:** Worked on scaling and improving the cloud data platform using Infrastructure as Code (IaC) and engineering best practices.

Used: Azure, Databricks, Terraform, Python, PySpark, Azure Devops, Docker, GIT, FOSS.

Data Science Portal | PostNL

- **Engineered Platform Infrastructure & CI/CD:** Built the core AWS platform infrastructure using AWS CDK (TypeScript) and implemented CI/CD pipelines for automated, streamlined deployments.
- **Developed Internal Developer Tooling:** Key creator of PostLit, an internal Python package integrating Streamlit with other internal platforms, automated PostNL branding to enable rapid PoC development for data scientists.
- **Accelerated Stakeholder PoCs:** Enabled data science teams to rapidly build, test, and showcase interactive web applications and proof-of-concepts for business stakeholders.

Used: AWS, AWS CDK, Python, TypeScript, Streamlit, Docker, GIT, OctoKit, FOSS.

Employment

2025 – Today

Data / AI / ML - Engineer | **Team Rockstars IT**

Amsterdam | Utrecht | Den Bosch | Arnhem

As a polyglot consultant and engineer, I had to jump into complex environments and delivering reliable results regardless of the existing stack. I adapt quickly to client needs—whether building scalable platforms (AI, Data, Applications) or core enterprise products.

Recent engagements include:

- **Rabobank** (Agentic Platform): Built a self-hostable agentic workflow platform using Temporal and Langfuse (IaC/DevOps), setting up initial agentic workflows, guardrails, and developed a structured way to handle agentic observability.
- **TenneT** (Data/ML Infrastructure): Contributed to an end-to-end legacy migration to Azure & Databricks, creating standardized DevOps templates and working on the Data platform infrastructure via IaC.

Driven by continuous learning and hands-on experimenting, I apply sound engineering practices to build

2022 – 2025

AI / Data / Cloud / Software - Engineer | **PostNL**

Den Haag

- Led development of the internal Data Science Portal, now hosting 40+ Data Scientists building Streamlit applications with 40+ apps in production. Built robust infrastructure using AWS Cloud Development Kit (CDK) with CI/CD pipelines, creating a central hub for innovation within the Data Science Community.
- As part of the AI Center of Excellence (CoE) at PostNL, established multiple generative AI projects using both proprietary and open-source models (HuggingFace). Helped develop PostNLChat, an internal ChatGPT-like interface, and created a generative AI workflow that automates phone conversation summarization for the sales team.

2022 – 2022

Junior Data Engineer | **Aurai**

Amsterdam

Upon completion of the traineeship, I transitioned into a consultant role, undertaking assignments that apply my data engineering expertise to solve client challenges.

Education

2023 – 2025

Universiteit van Amsterdam Msc.

Amsterdam

Master of Science: Information Studies: Data Science Track

2022 – 2022

Aurai – Data Engineering Traineeship

Amsterdam

Traineeship: Data Engineering

2018 – 2021

Universiteit van Amsterdam Bsc.

Amsterdam

Bachelor of Science: Informatiekunde | Minor Artificial Intelligence

Technical Stuff

Programming, scripting, and markup languages

- Python
- Typescript
- Shell
- SQL
- Pulumi
- Terraform
- YAML
- TOML
- Rust

(Web) Frameworks

- FastAPI
- Hono
- Flask
- Streamlit
- Marimo
- Astro
- React
- Vue
- Leptos

Platforms

- Linux
- MacOs
- Windows
- Azure
- AWS
- Docker
- Podman
- Kubernetes
- Databricks
- HuggingFace

Databases/Data-stores

- SQLite
- Postgres (+plugins)
- Milvus
- Chroma
- Redis
- SurrealDB
- MongoDB
- DuckDB

AI Technology, Tools & Frameworks

- OpenCode
- Github Copilot
- Claude
- Openrouter
- LiteLLM
- Langfuse
- Pydantic AI
- LangChain
- LangGraph
- LlamaIndex

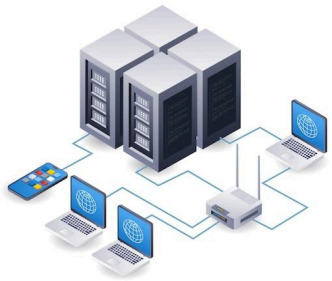

Tool & Libraries

- UV
- Poetry
- Ruff
- Ty
- Deno
- Bun
- Node
- Polars
- PySpark
- DataFusion
- Pandas
- IaC
- Pydantic
- Figma
- Tailscale
- Cookiecutter
- Copier
- Confluence
- Github Actions
- Azure Devops
- JFrog
- Nexus Cloud
- Zed
- VSCode
- PyTest
- MLFLow
- Grafana
- Excalidrwa
- DrawIO
- TaloOS

Projects

GUI OS AI Agent Benchmarking Framework

- Engineered Virtualized Testbed: Built a secure, QEMU-virtualized GUI desktop environment using Docker and Linux to enable safe, reproducible evaluation of autonomous system-level agents.
- Benchmarking Pipeline: Integrated Hugging Face's smolagents to automate performance testing for complex desktop interactions and multi-step OS workflows.
- Secure Execution: Automated containerized SSH and Jupyter interfaces to execute, log, and analyze agent system interactions safely isolated from host infrastructure.

Used: Python, Docker, HuggingFace, Smolagents, Linux, QEMU, SSH, Jupyter

melchiorhering/GUI-OS-AI-Agent-Benchmarking

Open Platform | Free & Open Source Software

An end-to-end, Kubernetes-native data platform bootstrapped via a modular Pulumi IaC stack for full-lifecycle deployment without architectural rebuilds. Integrates IDP-driven identity and access control, dynamic object storage, edge routing, and CDN integration. Operates on immutable OS baselines running containerized workloads with custom automation, backed by automated data orchestration, high-availability databases, telemetry, and native UI/API services.

Used: TypeScript, Python, Pulumi, Docker, Shell, Kubernetes, CI/CD, Self-Hosted.

/melchiorhering/open-platform